

Overview:

- Chapter 6 - Constraint Satisfaction
- Chapter 18 - Learning
- Chapter 7 - Logical Agents

Notes:

??? means do not have notes for this section, do not understand, or return to clarify.

Add comments for possible test question to be included in study guide

Chapter 6 - Constraint Satisfaction

- **Constraint Satisfaction Problems (CSP)**
 - Definition: solutions with caveats
 - Example: find a way to take classes such that I graduate in four years.
Constraints are prereq's, course availability, and funding
 - Previously, states were
 - Atomic - don't care about internal rep. except w/ respect to goal/heuristic
 - Mutated by actions to produce new atomic state
 - Now - **Factored representations**
 - States have **internal structure**
 - Structure can be **manipulated**
 - **Constraints** related different parts of the structure to one another and provide legal/illegal configurations
- CSP Definition
 - **Problem = {X, D, C}**
 - **X - set of variables** $X = \{X_1, X_2, \dots, X_n\}$
 - **D - set of domains** $D = \{D_1, D_2, \dots, D_n\}$
 - such that $X_i = x_i$ where x_i in D_i
 - **C - set of constraints** $C = \{C_1, C_2, \dots, C_n\}$
 - such that $C_i = \langle (C_a, C_b), \text{relationship}(C_a, C_b) \rangle$
- CSP Example: Map Coloring
 - Color territories on a map using 3 colors **such that no two colors are adjacent**
 - Can be represented as a graph with adjacent area connected by an edge
 - All variables have the same domain (the three colors)
 - **Constraint set consists of areas adjacent not being equal** *Resultant set*
- CSP Example: Scheduling
- **Constraint Types**
 - **Domain values** - time at which task begins $\{0, 1, 2, \dots\}$
 - **Precedence** constraints - $a + b \leq c$
 - Disjunctive constraints - e.g. can only do one thing at a time $\Rightarrow a + 10 \leq b$ **or** $b + 10 \leq a$
 - **Unary - single variable** ($z \leq 10$)
 - Binary - between **two variables** $z^2 > y$

- **Global** - constraints with 3+ variables can be **reduced to multiple binary/unary constraints**
- **Constraint graphs / hyper graphs** - see slides for picture
- **Binarization of constraints**
 - Convert n-ary constraints into **unary/binary** ones
 - Any arbitrary n-ary constraint can be converted to equivalent **unary** constraint
 - Example 1: Individual variables and their domains
 - **Variables:** $X = \{1,2\}$ $Y = \{3,4\}$ $Z = \{5,6\}$
 - **Encapsulated U** = $\{(1,3,5), (1,3,6), (1,4,5), (1,4,6), (2,3,5), (2,3,6), (2,4,5), (2,4,6)\}$ *all possibilities that are possible*
 - Introduce new encapsulated variable that is a Cartesian product of the domains of the individual variables! *All possibilities that satisfy throughout all possibilities*
 - Cartesian product $U = X \times Y \times Z$
 - This new encapsulated variable U contains all the unique combinations (8) *Let satisfy constraints*
 - Example 2: Original constraint and variables
 - **Constraint:** $X + Y = Z$
 - **Variables:** $X = \{1,2\}$ $Y = \{3,4\}$ $Z = \{5,6\}$
 - **Encapsulated U** = $\{(1,3,5), (1,3,6), (1,4,5), (1,4,6), (2,3,5), (2,3,6), (2,4,5), (2,4,6)\}$ *can get rid of some of*
 - **Encapsulated U (reduced)** = $\{(1,4,5), (2,3,5), (2,4,6)\}$
 - Create the encapsulated variable like before but perform a reduction based on the constraint! It will reduce the domain of U
 - <https://ktiml.mff.cuni.cz/~bartak/constraints/binary.html>
- **CSP Example: House Puzzle**
 - Row of houses each one has
 - color
 - Person with nationality
 - favorite candy
 - Favorite drink
 - Pet
 - All attributes distinct
 - Associate variables with a location
 - e.g. milk - 3 for house #3, cat - 4 for house 4
- **Implementing a CSP problem: Representation**
 - **Variables** - simple list
 - **Values** - mapping from variables to value lists e.g. python dictionary
 - **Neighbors** - mapping from variables to list of other variables that participate in constraints
 - **Binary constraints** - explicit value pairs, functions that return a boolean value
- **General Strategies for Solving CSP**
 - **Local consistency:** reduce set of possible values through constraint enforcement and propagation

- Perform search on **remaining possible states**
- **Node Consistency**
 - A variable is **node-consistent** if all values satisfy all unary constraints
 - Other unary conditions could further restrict the domain
- **Arc Consistency**
 - **Arc-consistent**
 - **Variable** - all binary constraints are satisfied for the variable
 - **Network** - if all variables in CSP are arc-consistent
 - Arc consistency only helps when some combinations of values preclude others
 - **3 Outcomes** (when all arcs are consistent)
 - One domain is empty => **no solution**
 - Each domain has a single value => **unique solution**
 - Some domains have more than one value => **may or may not be a solution**
 - In this case, arc consistency isn't enough to solve
- **Paths and k-consistency - higher levels of consistency, beyond our scope**
 - **Path consistency**
 - See if pair of variables consistent with 3rd variable, solved similarly to arc consistency
 - **K-consistency**
 - Given k-1 consistent variables, can we make Nth variable consistent (generalization of consistency)
- **Global Constraints**
 - Consider an “all different” constraint
 - Each variable has to have a distinct value
 - Assuming M variables and N distinct values
 - What happens when $M > N$? Impossible to solve, have to repeat a value
 - Extending this idea: ??? - review this
 - Find variables constrained to single value
 - Remove these variables and their values from all variables
 - Repeat until no variable is constrained to a single value
 - Constraints cannot be satisfied if
 - Variable remains with empty domain
 - There are more variables than remaining values
- **Resource contains (“at most”)**
 - Consistency checks - minimum values of domains satisfy constraint?
 - Domain restriction - are the largest values consistent with the minimum ones?
 - ???
- **Range Bounds**
 - Impractical to store large integer sets
 - Ranges can be used [min, max] instead
 - Bounds propagation can be used to restrict domains according to constraints
- Once all constraints propagated, search for solution

- **Naive Search**
 - Action picks a variable and a value.
 - N variables, domain size $D \rightarrow N \times D$ possible search nodes
 - Search on next variables
 - Backtrack when search fails
 - **Problems**
 - N variables with domain of size D
 - $N \times D$ choices for first, $(N-1)D$ for second $\rightarrow N!d^n$
 - Only d^n possible assignments
- **Modified Search**
 - CSPs are **commutative** \Rightarrow order of variable selection does not affect correctness
 - Each level of search handles specific variable
 - Levels have d choices, leaving us with d^n leaves
- **Backtracking Search**
- **Select-Unassigned-Variable**
 - Fail-first strategies:
 - Minimum remaining value heuristic:
 - Select the most constrained value; the one with the smallest domain
 - Rationale: probably most likely variable to fail
 - Degree heuristic:
 - Use the variable with the highest number of constraints on other unassigned variables
- **Order-Domain-Values**
 - Order of values within a domain **may or may not make a difference**
 - If goal is to produce all solutions or if there are no solution \Rightarrow order has no consequence
 - In other cases, we use a fail-last strategy and pick the value that reduces neighbors' domains as little as possible
- **Why fail-first for variable selection and fail-last for value selection?**
 - For variables, order does not matter. So we want to eliminate as many possible. For values, order can matter so if we eliminate a value, we might hurt ourselves later when we go to check it out.
- **Inference in Search \Rightarrow Forward Checking**
 - Check arc consistency with neighboring variables
 - Not needed if arc-consistency was performed prior to search
- **Maintaining Arc Consistency (MAC)**
 - Algorithm that propagates constraints beyond the node
 - AC3 algorithm with modified initial queue
 - Typical AC3 - all constraints
 - MAC - constraints between selected variable and its neighbors
- **Back-Jumping**

- Maintain a **conflict set** for each variable X :
 - A set of assignments that restricted values in X 's domain
- When conflict occurs, backtrack to last conflict added
- **Back-Jumping Implementation:**
 - On forward checks of X assigned to x ,
 - When X deletes a value from Y 's domain, add $X=x$ to Y 's conflict set
 - If Y is emptied, add Y 's conflict set to X 's and backjump
 - Easy to implement, build conflict set during forward check
 - However, what we prune is redundant to what we'd prune from forward checking or MAC searches
- More sophisticated backjumps - add notes here
- Conflict-Directed Back-Jumps - add notes for example?
- **Constraint-Learning and No-Goods**
 - Minimal set of assignments that caused problem - no-goods
 - Avoid running into problem by adding new constraint (or checking no-good cache)
- **Local Search CSPs**
 - Alternative to what we have seen so far
 - Assign everything at once
 - Search changes one variable at a time
 - Which variable? ???
- **Min-Conflicts Local Search**
 - Pretty effective for many problems, e.g. million queens problem can be solved in about 50 steps
 - Essentially **greedy search**, consequently:
 - **Local extrema**
 - **Can plateau**
 - Techniques discussed for **hill climbing** can be applied (e.g. simulated annealing, plateau search)
- **Structure of CSP Problems**
 - Improve search by **exploiting structure**
 - Independent **subproblems** - solve separately
 - Tree structure CSP
 - **Not on exam ???**
- **Tree Structure CSP**
 - Basic Ideas
 - Order variables (topological sort) such that constraints form a tree
 - Solve one variable at a time, propagate
 - Not all CSP constraints form trees
 - Transforming graphs with cycles into trees
 - Solve variable that reduces remaining conditions to a tree
 - Select a set of variables, a cutset that reduces problem to a tree after removal and examine with each possible assignment

■ Not on exam ???

Ch 18 - Learning

- Agents can learn to improve
 - **Inference from percepts**
 - Information about world evolution as the result of **changing world** or **action**
 - Utility estimators
 - Action choices either update condition-action maps or involve goal modification to maximize utility
- What we want to learn
 - **Mapping function**
 - Inputs are factored representation e.g. vector of values
 - **Outputs are**
 - Discrete (e.g. categorical)
 - Continuous
- **Type of Learning**
 - Inductive - learn map between input / output pairs
 - Deductive - creating rules that are logically entailed
- Learners vary based on feedback
 - **Unsupervised learning**
 - No explicit feedback
 - Goal is to cluster similar things
 - **Reinforcement learning**
 - Learner given rewards/punishments for actions
 - Ex: animal training w rewards
 - **Hybrids** are possible such as semi-supervised learning where a small set of labeled data accompanies large set of unlabeled data
- Caveat about labeled data sets
 - Labels referred to as "ground truth"
 - **Why be cautious with "ground truth"?** Error can be systematic / important?
- Supervised Learning
- How to choose amongst functions?
- Hypothesis spaces
- Supervised Learning
- **Decision Tree Learner**
 - Answers a series of **questions** to arrive at a solution
 - For now, restrict discussion to
 - Questions that have **categorical** (discrete answers)
 - **Binary** classification decisions
 - Decision trees -> good for justifying decision because you can see the decisions (white box)
- Constructing a tree from examples
- Quantity of information

- Amount of surprise that one sees when observing an event
 - If an event is rare, we can derive a large quantity of information (measured in bits) from it
- Expectation (review)
- Entropy
 - Defined as expected amount of information (average amount of surprise) and is usually denoted by the symbol H
- Example
- Restaurant Example
- Entropy and tree questions
- Tree Questions
 - Eq's have binary response
 - Suppose goal: separate mammals from birds
 - Question: does it fly?
- Tree Questions Entropy
 - Entropy is $E[I(P(X))] = E[-\log_2 P(X)]$
 - For binary categories, define short hand:
 - $Q = p / (p + n) \Rightarrow$ the positive rate
 - $1 - Q = n / (p + n) \Rightarrow$ the negative rate
 - $B(q) = E[I(P(X))] = -q \log_2 q - (1-q) \log_2 (1-q)$
 - Non-binary entropy calc:
- $$-\sum_{i \in \text{num_classes}} P(v_i) \log_2(v_i) == \sum_{i \in \text{num_classes}} P(v_i) \log_2(1/P(v_i))$$
 - Bird/mammal example
- Entropy and Tree Questions
- **Information Gain**
 - Goal: **reduce** the amount of information needed to represent the problem
 - We can represent the remaining entropy after dividing **data into d groups** with question **A** as follows:
 - **Remainder(A)** = some questions from slide

$$\text{Remainder}(A) = \sum_{k=1}^d \frac{p_k + n_k}{p + n} B\left(\frac{p_k}{p_k + n_k}\right)$$

- Expected value of entropy for all the categories
- And information gain as:
 - **Gain(A) = B(p/(p+n)) - Remainder(A)**
- **B = Binary Entropy**
- Information Gain Examples
- Decision Tree Learner
- Features and Overlearning
 - Useless features are not good for prediction, but a learner may pick up on random patterns in the training data and incorporate into rules

- Generalization and overfitting
 - Learning random patterns that don't affect function is called overfitting
 - For each leaf node, ask ourselves if good information gain
 - If node informative, keep it, otherwise discard
 - How do we know if our decisions were any good?
 - Goal was to **separate into positive and negative classes as well as possible**
 - Can we devise a statistic that lets us know if our observed split is statistically significant from the expected ratio? => Chi squared
- **Chi Squared Test**
 - Supposed decision tree splits a node into V categories
 - If node does not add any new info, we would expect # of class examples in each split roughly according to same proportion of examples
 - Let's restrict an example to
 - We can look at how much our categories differ from what would be expected if the proportion of categories did not change
 - When value is small, we are close to the original distribution
 - Test stat has distribution related to number of categories - 1 => **degrees of freedom**
 - Cumulative density function (CDF) -> integrate P(x dof) up to delta
 - The chi-squared test is used to determine **whether there is a significant difference** between the **expected frequencies** and the **observed frequencies** in one or more categories.
- **More Thoughts on Decision Trees**
 - Continuous/integer-valued attributes
 - Don't create infinite branches
 - Select split point
 - Sort values
 - Keep running total of number of +/- examples for each point in sorted list and pick separating point that give best separation
- Decision Tree Summary
 - Relatively straightforward learners
 - Recursively portion the feature space into hyperplanes
 - Sensitive to overtraining, have methods to prune
 - Easy for humans to understand
- Do I have a good hypothesis function?
 - Assume that are independent and identically distributed (iid)
 - ???
 - ???
 - We cross-validate the learner on a separate validation set
 - Problem: don't exploit all data
- K-fold cross validation
 - Extreme case: leave-one-out cross validation (aka jackknife)

- **Model selection**

- More complex models (e.g. more nodes in a decision tree) learn the training data better, but are they really better?
- Look at validation error

- **Loss**

- Loss functions are form of **utility function that provide cost for misclassification**
- Could be good to find whale and noted to misclassify nonwhale
- Some learners attempt to minimize loss
- Common loss functions ???

$$L_1(x, y, \hat{y}) = |y - \hat{y}| \quad \text{absolute loss function}$$

$$L_2(x, y, \hat{y}) = (y - \hat{y})^2 \quad \text{squared loss function}$$

$$L_{0/1}(x, y, \hat{y}) = \begin{cases} 0 & y = \hat{y} \\ 1 & \text{otherwise} \end{cases} \quad \text{0/1 loss function}$$

○

- **Generalization Loss**

- What is our loss when we use a novel data set e?
- The expected loss requires the distribution (X,Y) which we probably do not have:

$$GenLoss_L(h) = \sum_{(x,y) \in \mathcal{E}} L(x, y, h(x)) P(x, y)$$

○

- But we can estimate it empirically on a finite set of examples E of N samples:

$$EmpLoss_L(h) = \frac{1}{N} \sum_{(x,y) \in E} L(x, y, h(x))$$

○

○

○

- **Unreliability:** f may not be in H
- **Variance:** learners return different f's for different training sets
- **Noise:**
 - F may be noisy (e.g. stochastic component - different y's for the same x)
 - The training samples may have mis-measured attributes or incorrect labels
 - Might not have measured important attributes
- **Complexity:** learner may not achieve a global minimum

- **Regularization of Regression**

- ???

- **Linear Classification**

- Regression lines can be used to classify examples
- We look for a linear separating line (hyperplane when data is in R3 or higher)
- Suppose we pick a vector w perpendicular to the decision boundary
- Consider the dot product between w and an arbitrary point
- Once we have w , we can classify by taking the dot product and looking at the sign
- How do we choose w ?
 - Variant of the gradient descent rule used for regression is applicable
 - Recall regression rule

- **Perceptron Learning Rule**

- Iterate through data (1 iteration = 1 epoch) and repeat until no errors
- Convergence may be slow, but guaranteed for linearly separable data
- Most data sets not linearly separable
- When non linearly separable examples presented in random order, we will converge to a stable classifier if alpha decays linearly with epoch number

- **Learning with Logistic Regression**

- Remember update rule was:




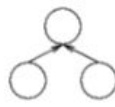
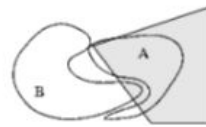

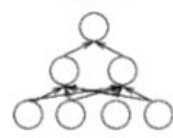
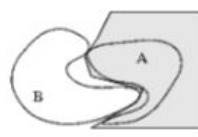
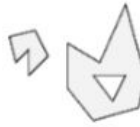
$$w_i = w_i - \alpha \frac{\partial}{\partial w_i} Loss(w)$$

- When the hard threshold is replaced with logistic regression:
 - Easier if we use labels $[0, 1]$ instead of $[-1, 1]$ as this range covered by function
 - Derivative of loss needs to be recomputed
 - $$w_i = w_i - \alpha (y - h_w(x)) h_w(x) (1 - h_w(x)) x_i$$
 - This results in smoother + more predictable learning curve

- **Connectionist Networks (Artificial Neural Networks)**

- Activation functions for perceptrons are **nonlinear**:
 - Hard threshold
 - Logistic regression (frequently called sigmoid function)
- **Linking perceptrons together** provides complex function modeling capability

- **Decision Boundary Capability as a Function of Network Depth**

Structure	Types of Decision Regions	Classes with Meshed Regions	Most General Region Shapes
One Layer 	Half Plane Bounded by Hyperplane		
Two Layer 	Convex Open or Closed Regions		
Three Layer 	Arbitrary (Complexity limited by number of nodes.)		

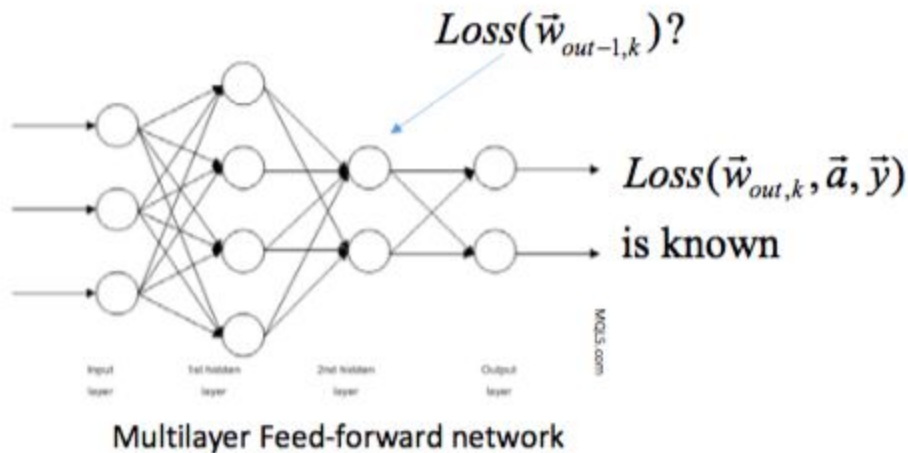
after Lipson 1992

- An Intuitive View of Neural Nets
 - Suppose we combine two perceptrons whose output functions are reversed
 - This could be used to model a ridge in output space
- Learning in a Neural Network
 - Similar to the regression problem, for output a and desired output y , we can find the loss gradient for each output node

$$\frac{\partial}{\partial w} \text{Loss}(w) = \frac{\partial}{\partial w} |y - h_w(x)|^2 = \frac{\partial}{\partial w} \sum_{k=1}^D (y_k - a_k)^2 = \sum_{k=1}^D \frac{\partial}{\partial w} (y_k - a_k)^2$$

$$a_k = \frac{1}{1 + e^{-w_{\text{output } k} \cdot \text{input}}}$$

- Use perceptron learning rule for the sum of the gradients at the output layer
- Back-Propagation
 - What should the targets be for the previous input layer?



-
- Back-propagating error (overview)
 - Error of the kth output: $\text{Err}_k = y_k - a_k$
 - We can compute gradient for any input node and apply regression rule
 - This gives new set of weights for output node
 - After applying update to the output layer, there still exists loss
 - We assign a portion of the loss to each of the input nodes based on their weight
 - This contribution is computed for each node of current layer
 - Look at sum of losses attributable to each node in previous layer -> sum of these provides us with loss to minimize
 - Repeat recursively
- ... gap in nodes ???
- **Neural Net Summary**
 - Supervised Learner
 - Training labels either
 - High value for class (n classes -> n output nodes)
 - Encoding of class information
 - Iterative training typically using a gradient descent algorithm (e.g. back propagation)
 - **Classification**
 - Present features to input nodes
 - Interpret output nodes for category
 - Disadvantages
 - Frequently hard to interpret
 - Many parameters require large data sets
 - Bad w imbalanced examples
 - Slow to train
 - Overfits easily regularization important
 - Advantages
 - Flexible, nonlinear learner

- Deep architectures very powerful
- **Non-Parametric Models**
 - Neural nets and decision trees have models with parameters
 - Decision node parameters: **attribute and cut-point/categories for sub-trees**
 - Neural nets: **weights and connections**
 - **Non-Parametric models:**
 - Cannot be characterized by **bounded set** of parameters
 - Simplest case: look at every example and use it to classify novel examples
 - Called **instance** or **memory-based learning**
- **Nearest Neighbors Models**
 - Use a distance metric to find the k closest neighbors, e.g. for continuous attributes
 - Use the plurality of labels that are the k closest
 - Good
 - Simplicity
 - Effective technique for low-dimensional data
 - Bad - searching is expensive with large training sets, but we can mitigate for this:
 - Trees - similar to decision tree (split on value, may at times need to search both sides)
 - Locally sensitive hash tables
 - Hash functions
 - Set of projections on to lines
 - Line projections discretized into buckets
 - Can be much more effective than tree approach
 - The Ugly
 - N points uniformly distribution in R^D unit hypercube
 - To capture $r = 0.01$, what edge length would we need in a random sample?
 - Samples are randomly distributed and total volume is 1, so we need a volume of r (0.01)
- **Support Vector Machines**
 - A **margin** is the distance to the closest examples on either side of a hyperplane
 - SVM approaches attempt to **maximize the margin**
 - Can only separate linear problems, but kernel function can project the data into a higher dimensional space where perhaps the data can be better separated
 - Maximal margins computed as functions of training examples
 - Summary
 - **SVM - non parametric technique**
 - In practice, **only small subset of training examples, the support vectors are required**
 - Training algorithm beyond our scope
- **Bias and Variance**

- Error in learning comes from two sources: **bias** and **variance**
 - **Bias** - large when learners make consistently incorrect predictions
 - **Variance** - large when different training sets result in different predictions
- **Ensemble Learning**
 - **Ensemble learners** are collections of weak learners that are combined to form robust classifier
 - **Weak learner** - simple learning algorithm that is likely to have high bias (e.g. **single node**, or stump, **of decision tree**)
 - Ensemble learners typically use collections of weak classifiers to reduce both bias and variance
- **Adaptive Boosting (ADABOOST)**
 - No notes ???

Chapter 7 - Logical Agents

- **Logical Agents**
 - Two key components:
 - **Representation** of the world
 - **Background information**
 - **Percepts**
 - **Ability to reason** : derive new information based on inference
- **Knowledge Base (KB)**
 - **Sentence** - statement about the agent's world
 - Based on **sensors**
 - Based on **background knowledge** (or possible learning)
- **Satisfaction and Entailment**
 - Suppose a is true in model M , then we state:
 - M satisfies a or equivalent
 - Or: M is a model of a
 - $M(a)$ means all models that satisfy a
 - Reason - entailment
 - $\alpha \models \beta$ if and only if $M(\alpha) \subseteq M(\beta)$
- **Examples of Entailment**
 - House is cornflower blue entails house is a shade of blue
 - $X = 0$ entails $xy = 0$
 - Second set must be bigger and contain it (basically subset)
- **Inference Algorithms**
 - Are sound if inference only devices entailed sentences
 - If our algorithm entailed $KB \models a_2$, we might fall into a pit!
 - Are complete if they can service any sentence that is entailed. Becomes an issue when the left-hand side is infinite
 - If KB is well grounded, our entailments should follow in the real world
- **Sentence Construction: Propositional Logic**

- Propositional logic is for the most part the logic you learned to program with
 - Sentence \rightarrow AtomicSentence | Complex Sentence
 - AtomicSentence \rightarrow true | false | literal
 - ComplexSentence \rightarrow (Sentence) | [Sentence]
 - | \sim Sentence \Rightarrow (negation)
 - | Sentence \wedge Sentence (conjunction)
 - | Sentence \vee Sentence (disjunction)
 - | Sentence \Rightarrow Sentence (implication)
 - | Sentence \Leftrightarrow Sentence (biconditional)
 - Operator priority is shown by order
- **Sentence Semantics**
 - Sentences reduced to true or false with respect to specific model
 - In the Wumpus cave
 - we might denote presence or absence of pit by literals indexed by location: $P_{x,y}$
 - Example: $P_{1,2}$, $P_{2,2}$, and $P_{3,1}$ that have true/false values in any given model
 - Models must specify values for each proposition
 - To resolve sentence: apply logical connectives to truth values
- **Knowledge Base Rules**
 - In the Wumpus save,
 - Denote pit & rumpus presence/absence y $P_{i,j}$ and $W_{i,j}$
 - There is a breeze if and only if there is a neighboring pit:
 - $B_{2,2} \Leftrightarrow P_{2,1} \vee P_{2,3} \vee P_{1,2} \vee P_{3,2}$
 - $B_{1,1} \Leftrightarrow P_{1,2} \vee P_{2,1}$
 - There is stench if the rumpus is in the next cavern
 - $S_{2,2} \Leftrightarrow W_{2,1} \vee W_{2,3} \vee W_{1,2} \vee W_{3,2}$
 - $S_{1,1} \Leftrightarrow W_{1,2} \vee W_{2,1}$
- **Percepts**
 - There is no pit at 1,1
 - $\sim P_{1,1}$
 - There is no breeze at 1,1
 - $\sim B_{1,1}$
 - There is a breeze at 2,1
 - $B_{2,1}$
- **Simple Entailment Algorithm**

```

def TruthTable-Entails(KB, a):
    symbols = proposition-symbols in KB and a
    return TT-Check-All(KB, a, symbols, {})

def TT-Check-All(KB, a, symbols, model):
    if empty(symbols):
        if pl-true(kb, model): # Does the KB entail the model?
            return pl-true(a, model) # Is entail the model?
        else return True # return true when KB does not hold

    else: # recursively enumerate the models
        (s, others) = (first(symbols), rest(symbols))
        return TT-Check-ALL(KB, a, rest, model U {s=True}) and
            TT-Check-ALL(KB, a, rest, model U {s=False})

```

- **Summary of Model Checking**

- Recursively generates all possible combinations of truth values for every symbol
- Checks if the knowledge base is a subset of a specific symbol value assignment
 - If so, returns if sentence a hold as well
 - Otherwise returns true as we don't care about whether a holds outside the KB (implies)
- Can we prove things without enumerating everything?

- **Concepts**

- **Logical equivalence**
- $a = B$ if they are true in the same set of models
 - Alternative definition:
 - Equivalent if they entail one another
 - $a \models B \iff B \models a$
- **Validity**
 - Sentences are valid (called **tautologies**) if they are true in all models
 - e.g. $(a \wedge b) \vee \neg a \vee \neg a$
 - **Deduction theorem**
 - For arbitrary sentences a and B, $a \models B$ iff $(a \Rightarrow B)$ is valid
 - So if we can show $(a \iff B)$ is a tautology, we know $a \models B$

- **Satisfiability**

- There exists **some model** such that sentence is true
- Sometimes easier to show something is valid by **showing its contradiction is not satisfiable**:
 - a is valid iff $\neg a$ is not satisfiable
 - if no model satisfies $\neg a$, then a must be true
- which leads to: $a \models B$ iff $(a \wedge \neg B)$ is not satisfiable
- Remember $a \models B$ iff $a \Rightarrow B \iff \neg a \vee B$

- **Inference Rules**

- ???

