

Introduction

Professor Marie Roch
Chapters 1 & 2, Russell & Norvig



Intelligence per Merriam Webster

- 1 a :** 1) the ability to learn or understand or to deal with new or trying situations : [REASON](#); *also* : the skilled use of reason
2) the ability to apply knowledge to manipulate one's environment or to think abstractly as measured by objective criteria (as tests)
b *Christian Science* : the basic eternal quality of divine Mind
c : mental acuteness : [SHREWDNESS](#)
- 2 a :** an [intelligent](#) entity; *especially* : [ANGEL](#)
b : intelligent minds or mind <cosmic *intelligence*>
- 3 :** the act of understanding : [COMPREHENSION](#)



What does it mean for a machine to be intelligent?

Turing Test – Can a human tell that they are interacting with a computer?



Variant of this competition occurs today with Loebner Prize (very restricted Turing test for chatbots)



Alan Turing 1912-1954
(see *The Imitation Game*
2014 historical for a
dramatization of his life)




Artificial intelligence

- Turing test - acting intelligent is sufficient
- Cognitive approach – Model the thinking of humans or other animals
- Laws of thought – Set of logic rules
drawback: difficult to model informal/general knowledge
- Rational agent – Logical inference is good, but not always applicable.



When will machines be intelligent?

(Armstrong, et al. 2014)



Eliza (Weizenbaum, 1966)

young woman: Men are all alike.

eliza: In what way?

young woman: They're always bugging us about something specific or other.

eliza: Can you think of a specific example?

young woman: Well, my boyfriend made me come here.

eliza: Your boyfriend made you come here?

young woman: He says I'm depressed much of the time.

eliza: I am sorry to hear you are depressed.

young woman: It's true. I'm unhappy.

eliza: Do you think coming here will help you not to be unhappy?

young woman: I need some help, that much seems certain.

eliza: What would it mean to you if you got some help?

young woman: Perhaps I could learn to get along with my mother.

Simple rules ---> disappointment

(X me Y) ---> (X you Y)

(I remember X) ---> (Why do you remember X just now?)

(My {family-member} is Y) ---> (Who else in your family is Y?)

(X {family-member} Y) ---> (Tell me more about your family)




Areas contributing to AI

- Philosophy
- Mathematics
- Neuroscience
- Psychology
- Linguistics
- Computer science

- Many application areas contribute as well (e.g. economics)



Intelligent agents



- Sensors provide perceptual input (percepts) of environment
- Agent makes decisions
- Actions carried out through actuators that may affect the environment



Intelligent agents

- Softbot – Software only agent
 - Available data are percepts
 - Examples: web-based reputation monitoring, game opponent



Task Environments

In what environment will the agent be operating?

- fully vs. partially observable
 - partially observable → uncertain state
- Rules are
 - known: Outcome (or outcome probabilities) are known
 - unknown – Outcomes must be learned



Task environments

- single- vs. multi- agent
- multiagent issues
 - cooperative vs. competitive
 - communication
 - randomization to prevent predictability



Task environments

- What happens when an agent acts?
 - deterministic – we know next state
 - stochastic
 - nondeterministic factors may influence (stochastic \rightarrow probabilities) leading to an *uncertain* state
- Decisions are
 - episodic – Next decision only depends on state
 - sequential – Next decision dependent on previous ones



Task environments

- State can be
 - static – does not change while agent is deciding next action
 - dynamic – Environment constantly changing
 - semidynamic – Environment static, but performance is time dependent



See p. 45 Figure 2.6 for example task environments

Quick and dirty Python 3.x

- About the language
 - Interpreted high level language
 - Reasonably simple to learn
 - Rich set of libraries
- For details, see texts in syllabus or www.learnpython.org or www.diveintopython.net
- Python comment
comment from hash character to end of line



Python data types

- Numbers: 42.8 or 9
- Strings: single or double quote delimited
 'hi there' "Four score and seven years ago..."
- Dictionaries: Python's hash table
 quotes = dict() # new dictionary
 quotes["Lincoln"] = "Four score and seven years ago..."
 OR
 quotes = {"Lincoln" : "Four...",
 "Roosevelt": "The only thing we have to fear..."}



Python data types

- Sequences
 - Lists ["Four", "score", "and"]
 - tuples ("Four", "score", "and")
- Difference between tuple and list
 - List – can grow or shrink
 - Tuple – Fixed number of elements
 - Faster
 - Can be used as hash table indices



Python data types

- None – special type for null object
- Booleans: True, False
- Variables are untyped



Python Expressions

- assignment: `count = 0`
- list membership: `value in [4, 3, 2, 1]`
- indexing: `listvar[4]`, `tuplevar[2]`
- slices:
 - `listvar[0:N]` → items 0 to N-1
 - `listvar[:N]` → items 0 to N-1
 - `listvar[3:]` → items 3 to end
 - `listvar[0:2:N]` → even items at 0, 2, 4, ...
 - `listvar[-4:-1]` → 4th to the last to 2nd to the last
- logical operators: `and`, `or`, `not`



Python expressions

- comparison operators: `<` `>` `>=` `<=` `!=`
- basic math operators: `+` `-` `/` `*`
- exponentiation: `x ** 3` # x cubed
- bitwise operators: `&` `|` `~` and `^` (xor)



Python control structures

- Use indentation to denote blocks
- Conditional execution

```
if expression:  
    statement(s)  
elif expression:  
    statements(s)  
else:  
    statement(s)
```



Python control structure

- Iteration

```
done = False  
while not done:  
    statements(s)  
    done = expression  
  
for x in range(10): # 0 to 9  
    print(x)  
    print("x={}".format(x))
```

alter iteration with break and continue



Python functions

```
def foobar(formal1, formal2, formal3=None):  
    "foobar doesn't do much" # doc string  
    statement(s)  
    return value
```

- formal3 defaults to None if not supplied
- Variable scope rules
 local, enclosing function, global, builtin names



Python objects

```
class Board:  
    "Grid board class"  
    def __init__(self, rows, cols): # constructor  
        "construct a board with specified rows and cols"  
        self.rows = rows  
        self.cols = cols  
        # list comprehension example  
        self.board = [[None for c in range(cols)] for r in range(rows)]  
    def place(self, row, col, item):  
        "place an item at position row, col"  
        self.board[row][col] = item  
    def get(self, row, col):  
        "get an item from position row, col"  
        return self.board[row][col]
```



Python objects

- Create: `b = Board(8,8)`
- `b.place(2, 7, 'black-king')`
- `b.get(2,7)`
 "black-king"



Python

- Integrated development environments
 - Eclipse with PyDev ← what I use
 - Komodo (ActiveState)
 - Pycharm
 - others (see Python.org)
- Versions of Python
 - Python.org – stock Python
 - PyCharm
 - Anaconda – bundles with lots of libraries and and ID



Agent structure

- An agent's architecture consists of
 - data structures
 - code
- Simplest agent: table driven

function table-driven-agent(percept) returns action

 persistent: percepts (sequence, empty at first)

 table of actions indexed by
 percept sequence

 percepts.append(percept)

 return lookup(percepts, table)

What's wrong here?



Simple-reflex agents

- Ignores percept history, uses the current one
- Productions (aka conditions-action) decide action, e.g.
 - person waving → wave
 - person smiling → smile
 - person swinging hammer towards me → duck!



Model-based reflex agents

- Add internal state
- New percepts update the state
- Productions based on percept and state



Goal-based agents

- Agent works to achieve a specific state
- Usually requires: Search and Planning



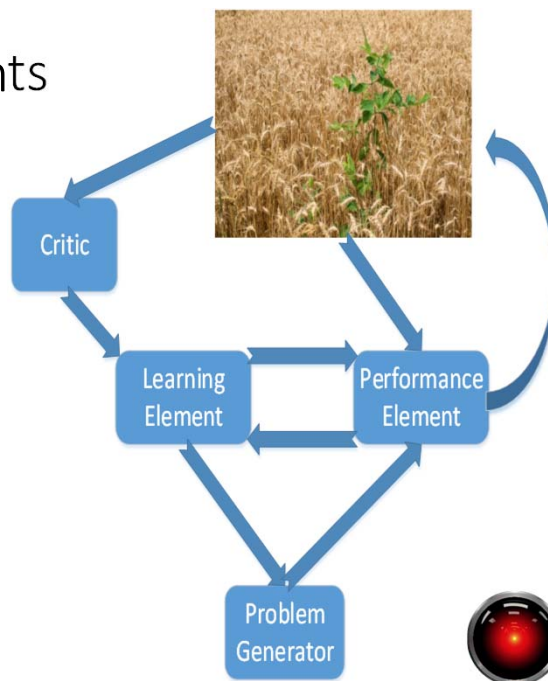
Utility-based

- Based on utility theory: The idea of how useful or happy something makes you.
- Decisions are made to maximize the expected utility.



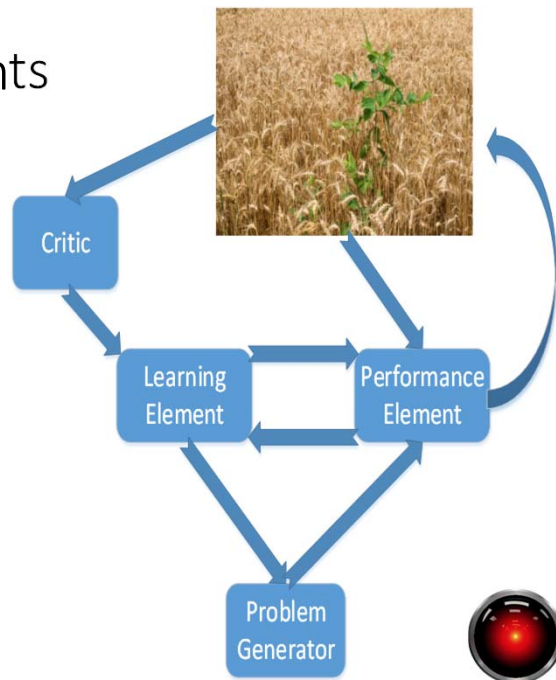
Learning agents

- Decision rules are adjusted based on performance over time



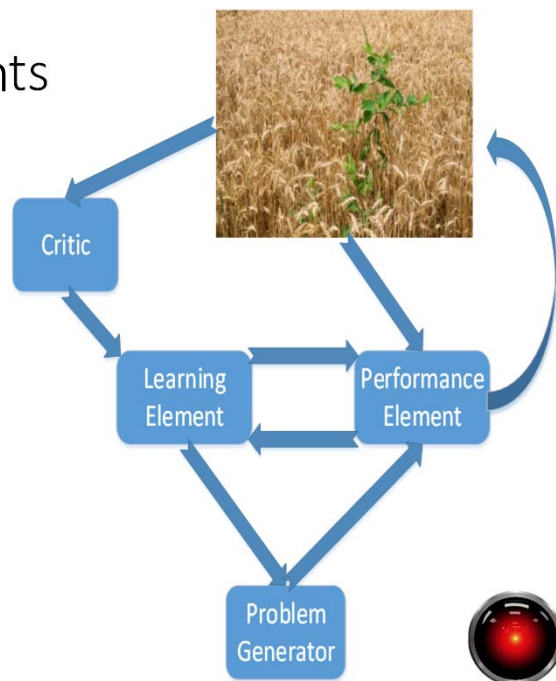
Learning agents

- Learning element uses knowledge from performance element to modify rules



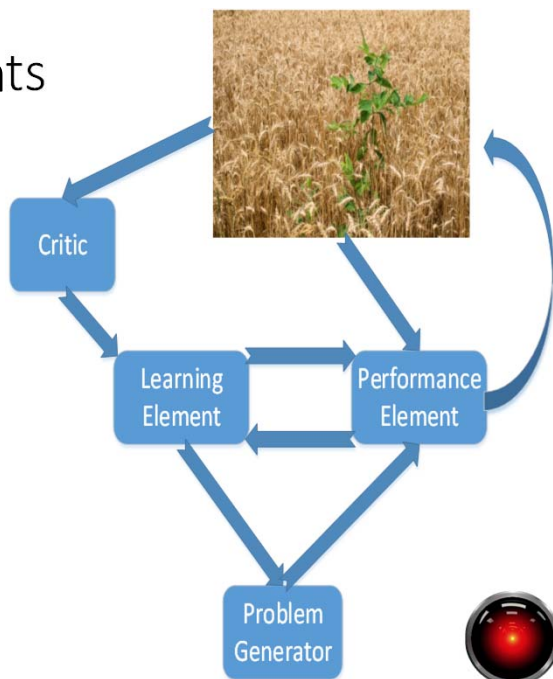
Learning agents

- Critic monitors the environment to see how the system is doing and keeps learning element from going off track.



Learning agents

- Critic monitors the environment
 - Provides utility of the current percepts
 - Allows learning element to learn useful goals



Learning agents

- Problem generator
 - Uses learning element's goals to suggest experiments
 - Experiments may lead to the learning element improving the performance element

