

Search

Professor Marie Roch
Chapter 3, Russell & Norvig



Solving problems through search

- State – atomic representation of world
- Goal formulation
 - What objective(s) are we trying to meet?
 - Can be represented as a set of states that meet objectives: goal states
- Problem formulation
 - Decide actions and states to reach a goal



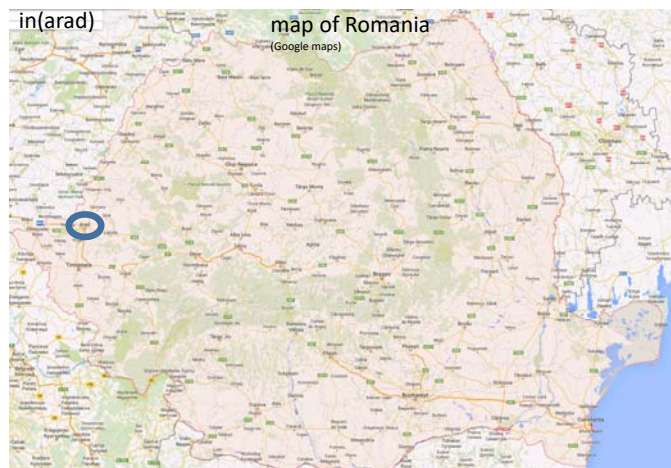
Search

- Assume environment is
 - observable
 - discrete (finite # of actions)
 - deterministic actions
- Search process returns a plan:
set of states & actions to reach a goal state
- Plan can be executed




Search problem components

- Initial state



Search problem components

- Initial state
- Actions
 - function that returns set of possible decisions from a given state
 - $\text{actions}(\text{in}(\text{arad})) \rightarrow \{\text{go}(\text{sibiu}), \text{go}(\text{timisoara}), \text{go}(\text{zerind})\}$



Abstract view of Romanian roads
(Russel and Norvig 2010, Fig 3.2)

Note: Abstractions are valid when we can map them onto a more detailed world



Search problem components

- Initial state
- Cost
 - Each action has a step cost:
 $\text{cost}(\text{in}(\text{arad}), \text{go}(\text{zerind}), \text{in}(\text{zerind})) = 75$
 - A path has a cost which is the sum of its step costs:

path: $\text{in}(\text{arad}), \text{in}(\text{zerind}), \text{in}(\text{Oradea})$
has cost: 146 (75+71)



Abstract view of Romanian roads
(Russel and Norvig 2010, Fig 3.2)



Search problem components

- Initial state
- Actions
- Cost
- Transition model
Function that reports the result of an action applied to a state:
`result(in(arad),go(zerind) → in(zerind)`



Arad




Zerind



Search problem components

- Initial state
- Actions
- Cost
- Transition model
- Goal predicate
Is the new state a member of the goal set?
goal: {in(bucharest)}




Any path that reaches a goal is a *solution*, the lowest cost path is an *optimal solution*.




Sample toy problem

- n-puzzle



8-puzzle and one possible goal state
[Figure 3.4 R&N 2010]

- n-queens



8-queens state
[Figure 3.5 R&N 2010]

see text for other examples



Constructing a problem: n-queens

- States

1. complete-state:
 - n-queens on board
 - move until no queen can capture another.
2. Incrementally place queens
 - initial empty board
 - add one queen at a time



Incremental n-queens

- state: Any arrangement of $[0, n]$ queens
- initial state: empty board
- actions: add queen to empty square
- transition model: new state with additional queen
- goal test: n queens on board, none can attack one another



Incremental n-queens

- A well-designed problem restricts the state space

- Naïve 8 queens
 - 1st queen has 64 possibilities
 - 2nd queen has 63 possibilities...

$$64 \times 63 \times 62 \dots \approx 1.8 \times 10^{14}$$

- Smarter:
 - Actions only returns positions that would not result in capture
 - State space reduced to 2057 states.



designed by
Christine Kawasaki-Chen

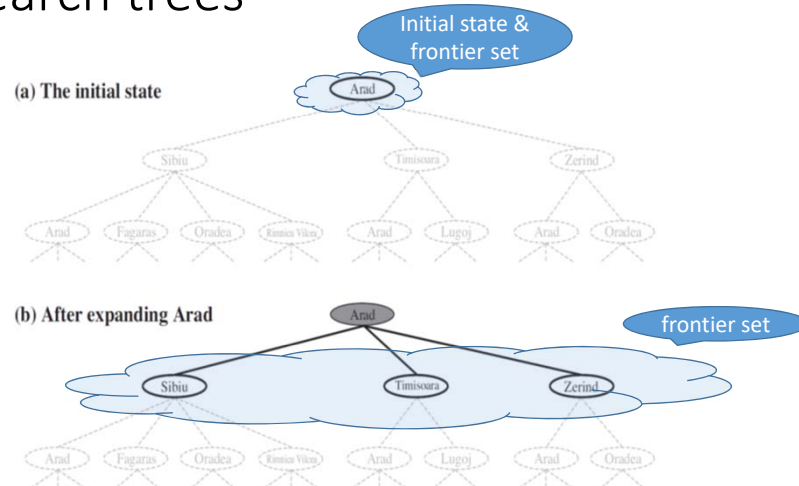


Classic real-world problems

- route-finding problem
 - transportation (car, air, train, boat, ...)
 - networks
 - operations planning
- touring problem
 - Visit a set of states ≥ 1 time
- traveling salesperson
 - Visit a set of state 1 time
- Others: VLSI layout, autonomous vehicle navigation & planning, assembly sequencing, pharmaceutical discovery



Search trees

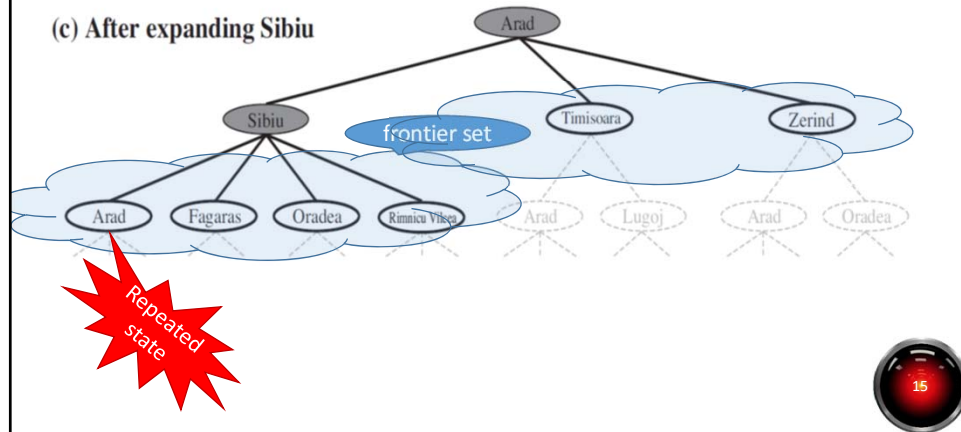


[Figure 3.6 R&N 2010]

frontier set also known as an open list



Search trees




Search tree

- Frontier set consists of leaf nodes
- Redundant paths occur when
 - \exists more than 1 path between a pair of states
 - cycles in the search tree (loops) are a special case



Redundant paths

"Those who cannot remember the past are condemned to repeat it"



George Santayana,
Spanish-American philosopher 1863-1952

- Sometimes, we can define our problem to avoid cycles
e.g. n-queens: queen must be placed in the leftmost empty column
- Otherwise: Explored set
 - Track states that have been investigated
 - Don't add any actions that have already occurred



Tree Search

```
function tree-search(problem)
  frontier = problem.initial_state()
  done = found = False
  while not done
    node = frontier.get_node() # remove state
    if node in problem.goals()
      found = done = True
    else
      frontier.add_nodes(results from actions(node))
      done = frontier.is_empty()
  return solution if found else return failure
```



Graph Search

```
function graph-search(problem)
    frontier = problem.initial_state()
    done = found = False
    explored = {} # keep track of nodes we have checked
    while not done
        node = frontier.get_node() # remove state
        explored = union(explored, node)
        if node in problem.goals()
            found = done = True
        else
            # only add novel results from the current node
            nodes = setdiff(results from actions(node), union(frontier, explored))
            frontier.add_nodes(nodes)
            done = frontier.is_empty()
    return solution if found else return failure
```



Search architecture

- Node representation
 - state
 - parent – ancestor in tree
 - allows us to find the solution from a goal node by chasing pointers and reversing the path
 - action – Which action was used on parent to generate this node
 - path-cost – What is the cost to reach this node from the tree's root. Usually denoted $g(n)$.



Search architecture

```
function child-node(problem, node, action)
  child.state = problem.result(node.state, action)
  child.parent = node
  child.path_cost = node.path_cost +
    problem.cost(node.state, action, child.state)
  return child
```

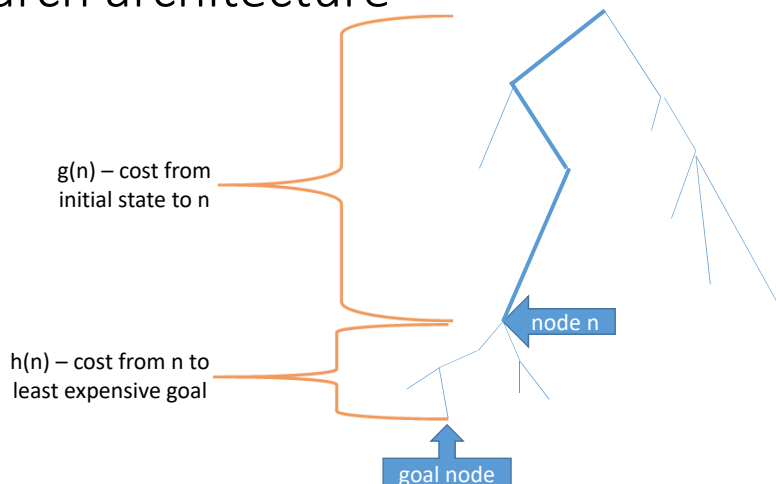


Search architecture

- frontier set is usually implemented as a queue
 - FIFO – traditional queue
 - LIFO – stack
 - priorityWe will develop a way such that it can always be a priority queue.
- Explored set – Need to make states easily comparable
 - hash the state or
 - store in canonical form (e.g. sort visited cities for traveling salesperson problem)



Search architecture



$g(n)$ and $h(n)$ are frequently not known precisely.
Estimates are denoted or $g'(n)$ & $h'(n)$ or $\hat{g}(n)$ & $\hat{h}(n)$



A generic graph search algorithm

```
function tree-search(problem)
    frontier = problem.initial_state() # priority queue based on lowest cost
    done = found = False
    explored = {} # keep track of nodes we have checked
    while not done
        node = frontier.get_node() # remove state
        explored = union(explored, node)
        if node in problem.goals()
            found = done = True
        else
            # only add novel results from the current node
            nodes = setdiff(results from actions(node), union(frontier, explored))
            for n in nodes
                estimate a cost  $g'(n) + h'(n)$ 
            frontier.add_nodes(nodes) # merge new nodes in by estimated cost
            done = frontier.is_empty()
    return solution if found else return failure
```






Uninformed (blind) search

- No awareness of whether or not a state is promising
- Strategies depend on order of node expansion
 - breadth-first
 - uniform-cost
 - depth-first
 - variants: depth-limited, iterative deepening, bidirectional
- Note: Text uses different queue types for frontier, with generic search algorithm everything is a priority queue, smallest values first.



Breadth-first search

- $\forall n \ g'(n) = \text{depth}(n)$ and $h'(n) = k$ (e.g. $k=0$)



Breadth-first search

- Guarantees
 - completeness – will find a solution
 - best path *if cost is a nondecreasing $f(\text{depth})$*
- How can we measure performance?
 - Time complexity
 - Space complexity



Complexity (CS 310 material)

- Measure of the number of operations (time) or memory (space) required
- Analysis of performance as the number of items n grows:
 - worst case
 - average case
- Example:

There are $T(n)=4n^2+1$
arithmetic operations

```
def foobar(n):
    x = 0
    for i in xrange(n):
        for j in xrange(n):
            x = x + i*i + j*j
    return x * x
```



Complexity

- We define “big oh” of n as follows:

$$T(n) \text{ is } O(f(n)) \text{ if } T(n) \leq kf(n) \\ \text{for some } k \text{ \& } \forall n > n_0$$

- Role of k and n_0

Coefficients of highest order polynomial aren't relevant.

- Implications:

- $T(n) = 4n^2 + 1 \rightarrow O(n^2)$
- $T'(n) = 500n + 8 \rightarrow O(n)$

For some small values of n , $T(n)$ is better, but as n increases $T(n)$ will be worse. Using the big-oh notation abstracts this away and we know in general that the second algorithm is better.



Search complexity


Measured with respect to search tree:

- Complexity is a function of
 - Branch factor – max # of successors
 - Depth of the shallowest goal node
 - Maximum length of a state-space path
- Time measurement: # nodes expanded
- Space measurement: maximum # nodes in memory




Search complexity

- “Search cost” – time complexity
- “Total cost” – time and space complexity
Problematic to fuse the metrics...



Breadth-first search performance

- Assume branch factor b
- Time complexity:
$$b + b^2 + b^3 + \dots + b^d = O(b^d) *$$
- Space complexity
 - Every generated node remains in memory, $O(b^{d-1})$ in explored and $O(b^d)$ in frontier.

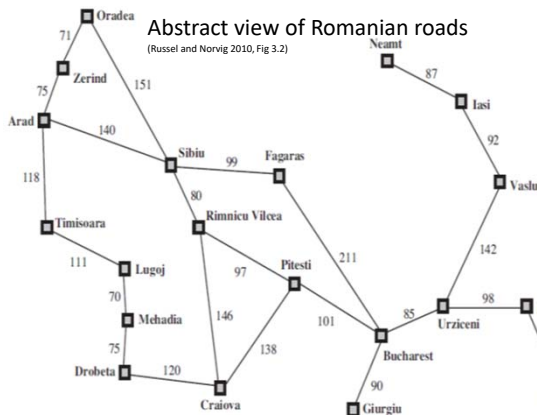


* See text for discussion of $O(b^d)$ vs. $O(b^{d+1})$



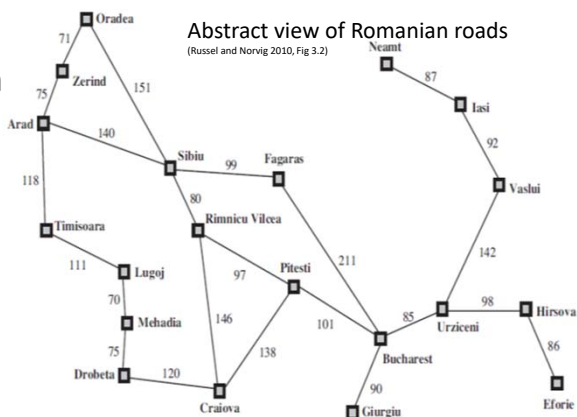
Uniform-cost search

- Similar to breadth-first, $g'(n)$ uses edge costs
- $\forall n, g'(n) = g(n)$ and $h'(n) = k$
- Nodes are expanded in order of optimal cost \rightarrow optimal solution
- Complexity function of minimum cost for all actions




Depth-first search

- Deepest node is expanded first
- $\forall n, g'(n) = k$ and $h'(n) = -\text{depth}(k)$
- Non-optimal
- Incomplete search
- Why bother?



Depth-first search (DFS)

- DFS will explore other paths when there are no successors.
- Fast! If you hit the right path... but the average case analysis is $O(b^m)$ where m is maximum depth.
- Space complexity is better: $O(bm)$



Iterative deepening

- Prevents infinite loops of depth-first search
- Basic idea
 - Depth-first search with a maximum depth
 - If the search fails, repeat with a deeper depth



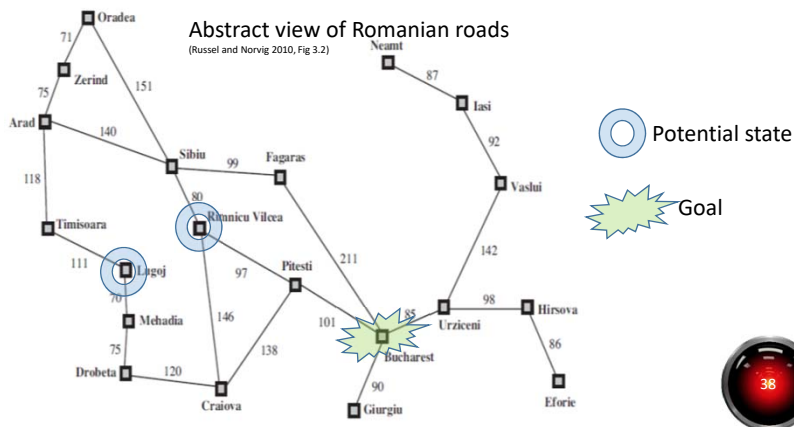
Uninformed search

- Other variants exist
- For large search spaces, it is generally a bad idea



Informed, or heuristic, search

- General idea: Can we guess the cost to a goal based on the current state?



Heuristic

- $h(n)$ – Actual cost from a search graph node to a goal state *along the cheapest path*.
- $h'(n)$ – An estimate of $h(n)$, known as a heuristic.

Note that your text does not make a notational distinction between the actual cost and the estimated one and always uses $h(n)$, so we will frequently follow suit.



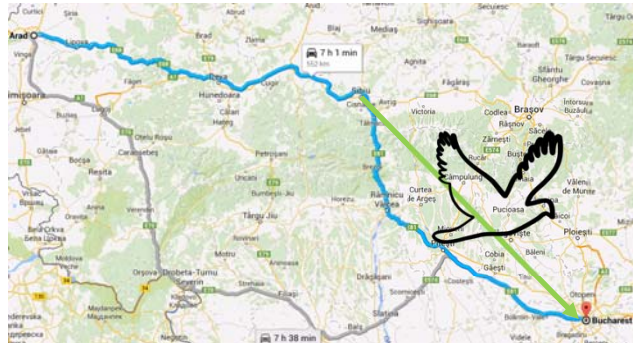
Heuristic

- $h(n)$ is always ≥ 0
- $h(n)$ is problem specific
- Estimators of $h(n)$ are similar.
- One can think of a heuristic as an educated guess.
We will look at how to construct these later...



Greedy best-first search

- $g(n) = 0$, $h(n)$ is heuristic value
- Example $h(n)$ for Romania example:
as the bird flies distance



A* Search


- “A-star” search uses:
 - $g(n)$ = cost incurred to n
 - $h(n)$ = estimate to goal
- A^* is the estimated cost form start to goal through state n



Heuristic properties

- **admissible** – $h'(n)$ is admissible if it *never overestimates* the cost to goal


One can think of it as optimistic: $h'(n)$



- **consistency** (aka monotonicity) – $h'(n)$ is consistent if

consistency:

$$h'(n) \leq \text{cost}(n, \text{action}, n') + h'(n')$$



Note: We are being careful about distinguishing the heuristic estimator $h'(n)$ from the actual distance $h(n)$

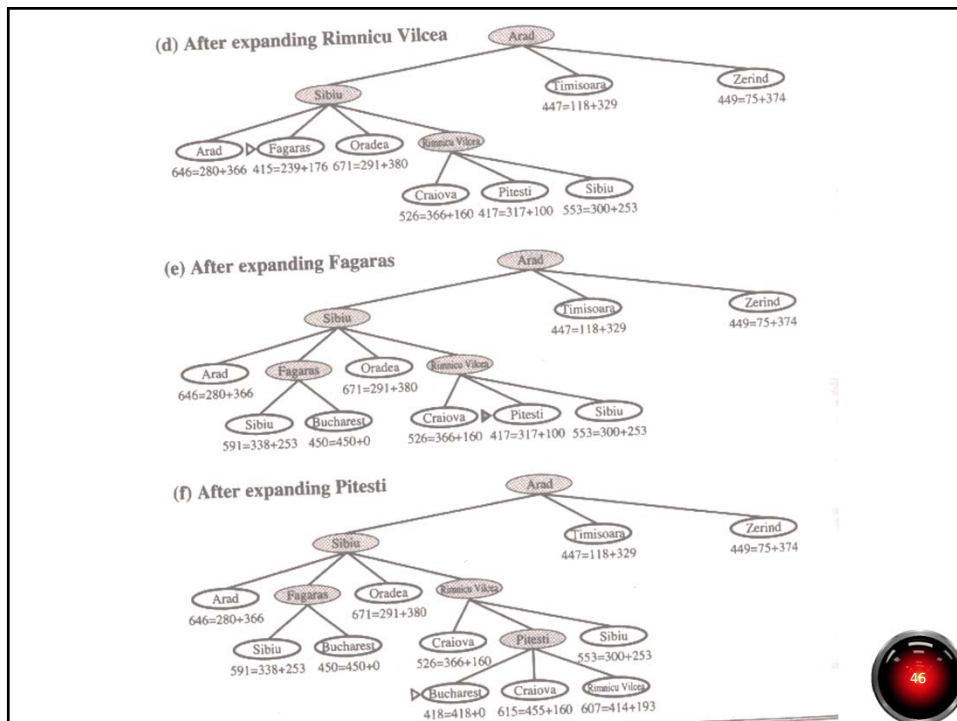
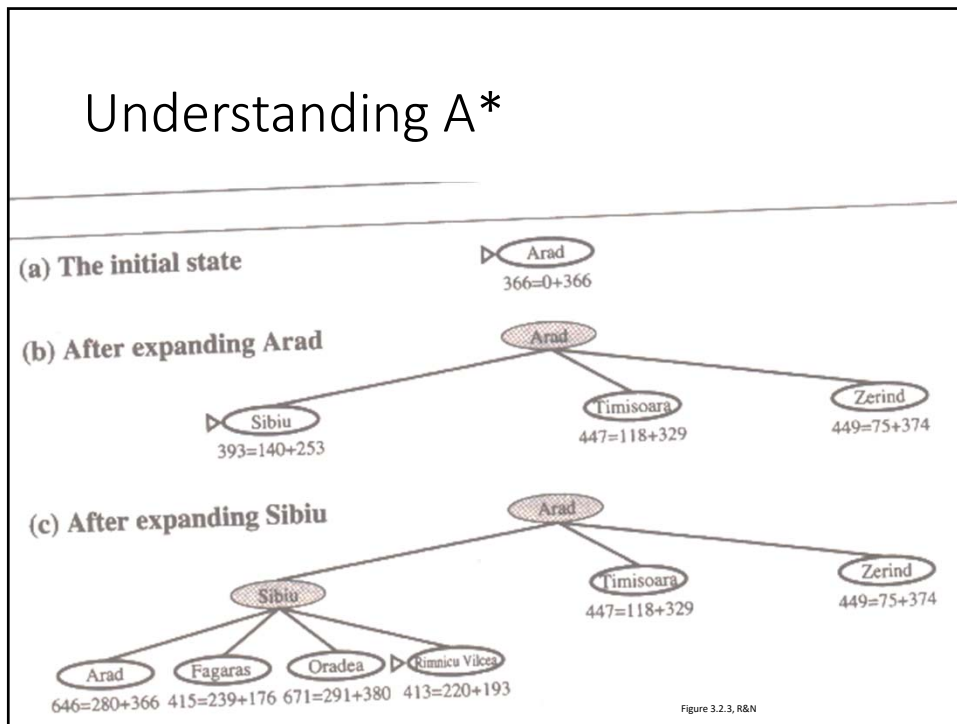


Heuristic properties

- Every consistent heuristic is also admissible.
- A* is guaranteed to be:
 - for trees
A* optimal if $h'(n)$ is admissible
 - for graphs
A* optimal if $h'(n)$ is consistent




Understanding A*




Understanding A* optimality

Consistency revisited:
the \blacktriangle inequality – the
sum of any two sides \geq
third side



$$h'(n) \leq c(n, \text{action}, n') + h'(n')$$

If h' consistent and costs
are nonnegative, values of
 $f(n)$ along any path are
nondecreasing.



Understanding A* optimality

- Suppose we pick node n
- Is the path to node n 's state optimal?

Proof by contradiction

Suppose a better path to the same state is in node b .

As b and n have the same state, so $h(n) = h(b)$.

Relative position in queue will be driven by $g(n)$ and $g(b)$.

If b is better, $g(b) < g(n)$ and we would have picked b first



Understanding A* optimality

- When $h(n)$ is consistent, the properties of:
 - nondecreasing values of $f(n)$
 - guarantee that we pick the best path to nensure that the first goal node we find is optimal.
- Completeness holds when there are a finite number of nodes with $f(n) < \text{the optimal cost}$



Limitations of A*

- Need to find a heuristic
- Show it is consistent (for graph search) if optimal goal is required.
- Show the graph is finite for nodes with cost lower than the optimal one if completeness is required
- Note: expanded set requires nodes in memory and is a frequent limitation of A*



A* variants

- iterative deepening A*
Same idea as iterative depth-first search, but we limit on $f(n)$
- SMA* - simplified memory A*
 - When memory is full
 - drops worst frontier node (highest $f(n)$)
 - stores that value in parent, and will only reconsider branch when everything looks worse than the stored value
 - Details beyond our scope



Heuristic search summary

- A* can still have problems with space complexity
 - iterative deepening A*
 - other alternatives listed in text
- Complexity of A* search is tricky, but is related to
 - the error in the heuristic, $h(n) - h'(n)$
 - and solution depth



Developing heuristics

- Requires
 - knowledge of problem domain
 - thinking a bit (usually)
- Effort to show that heuristic is
 - admissible
 - consistent

7	2	4
5		6
8	3	1

Start State

- What heuristics could we use for the N-puzzle



N-puzzle heuristics

- Common heuristics
 - $h_1(n)$ – Number of misplaced tiles
 - $h_2(n)$ – Sum of Manhattan¹ distance of tiles to solution
- Are these
 - admissible? (never overestimates)
 - consistent? (non-decreasing path cost)

8		6
5	4	7
2	3	1

¹ Also known as city-block distance, the sum of vertical and horizontal displacement.



Heuristics and performance

- Branching factor
 - Measured against a complete tree of solution depth d
 - Suppose A^* finds a solution at
 - depth 5
 - 52 nodes expanded (53 with root)
 - A complete tree of depth 5 would have

$$52 + 1 = b^* + (b^*)^2 + (b^*)^3 + (b^*)^4 + (b^*)^5$$
 where b^* is the branch factor
 - Using a root finder for $1(b^*)^5 + 1(b^*)^4 + 1(b^*)^3 + 1(b^*)^2 + 1(b^*)^1 - 53(b^*)^0 = 0$
we see $b^* \approx 1.92$



Heuristics and performance

- 8-puzzle example averaged over 100 instances

d	Search Cost (nodes generated)			Effective Branching Factor		
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	–	539	113	–	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

- branch factors closer to one are better



Finding heuristics

- Okay, developing a heuristic is hard
- Can we make it easier?




Relaxed problem heuristics

- Let's return to the N-puzzle
- Suppose we allowed
 - A tile to move onto the next square regardless of whether or not it was empty.
 - A tile to move anywhere.
- These are relaxations of the rules



Relaxed problems

We can think of these as expanding the state space graph.



Relaxed problem heuristics

- The original state space is a subgraph of the new one.
- Heuristics on relaxed state space
 - Frequently easier to develop
 - If admissible/consistent properties hold in relaxed space, they also hold in the problem state space.




Relaxation

- Can specify problem in a formal language, e.g.
 - move(A,B) if
 - verticalAdjacent(A,B) and
 - horizontalAdjacent(A,B) and
 - isempty(B)
- Possible relaxations
 - move(A,B) if adjacent(A,B)
 - move(A,B) if isempty(B)
 - move(A,B)



Automatically generated heuristics

With a formal specification of the problem there exist algorithms to find heuristics (beyond our scope, e.g. ABSOLVER)



Multiple heuristics

- Regardless of how generated, one may develop multiple heuristics for a problem
- We can merge them


$$h'(n) = \max(h'_1(n), h'_2(n), \dots, h'_i(n))$$

why maximum?



Pattern database heuristics

- Can we solve a subproblem?



- If we can, we can store its $h(n)$



Pattern database heuristics

- Cost usually found by searching back from goal nodes.
- Worth it if the search will be executed many times.
- Sometimes patterns are disjoint. If so, the heuristic costs may be added (doesn't work for 8 puzzle)



Learning heuristics

- Use experience to learn heuristics
- Beyond our reach for now... (machine learning)



Heuristic summary

(rough outline, no substitute for a little thought)

